

Relationship between fixed effects, random effects, GLS, and penalized regression; confounding; model misspecification

```
In [1]: library(data.table)
library(ggplot2)
library(lme4)
library(parallel)
library(penalized)
library(xtable)

options(repr.plot.width=6, repr.plot.height=4)
theme_set(theme_bw())
```

Loading required package: Matrix
Loading required package: survival
Welcome to penalized. For extended examples, see vignette("penalized").

Prepare Data

```
In [2]: literacy.covariates.data <- data.table(subject=c(1:91))
literacy.covariates.data[,base.age:=1 + 0.1*(subject - 1)]
literacy.covariates.data[,`0`:=base.age]
literacy.covariates.data[,`1`:=`0` + (1 + (10 - base.age)/10)^2]
literacy.covariates.data[,`2`:=`1` + (1 + (10 - base.age)/10)^2]
literacy.covariates.data[,subject:=factor(subject)]
literacy.covariates.data <- melt(literacy.covariates.data, id.vars=c('subject', 'base.age'),
                                value.name='total.age', variable.name='follow.per')
literacy.covariates.data[,follow.per:=factor(as.numeric(follow.per) - 1, ordered=TRUE)]
setkey(literacy.covariates.data, subject, follow.per)
literacy.covariates.data[,delta.age:=total.age - base.age]
head(literacy.covariates.data)
```

subject	base.age	follow.per	total.age	delta.age
1	1.0	0	1.0000	0.0000
1	1.0	1	4.6100	3.6100
1	1.0	2	8.2200	7.2200
2	1.1	0	1.1000	0.0000
2	1.1	1	4.6721	3.5721
2	1.1	2	8.2442	7.1442

Response

```
In [3]: make.base.literacy <- function(base.age) {
  (10 - base.age)^2
}

make.response <- function(covariates, variance=100) {
  make.base.literacy(covariates$base.age) +
  covariates$delta.age +
  rnorm(n=nrow(covariates), sd=sqrt(variance))
}
```

Visualization

```

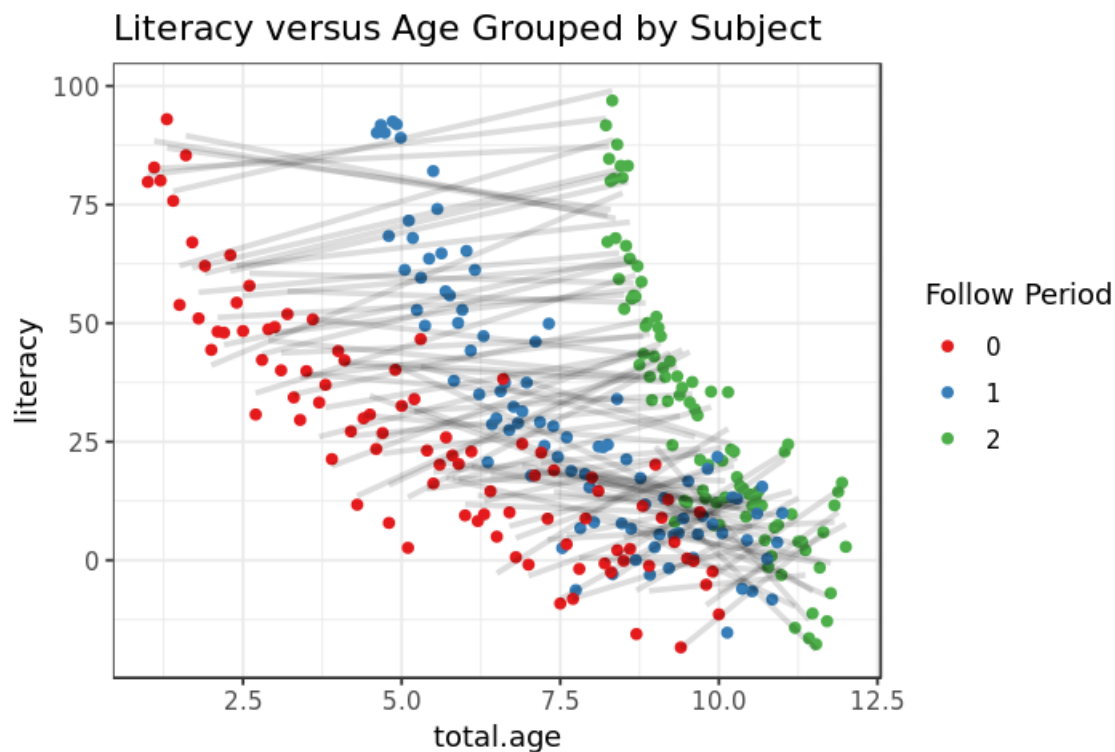
In [4]: set.seed(2021)
literacy.data <- copy(literacy.covariates.data)
literacy.data[,literacy:=make.response(literacy.data)]

p <- ggplot(literacy.data, aes(x=total.age, y=literacy)) +
  geom_point(aes(color=follow.per)) +
  scale_color_brewer('Follow Period', palette='Set1') +
  geom_smooth(aes(group=subject), method='lm', se=FALSE, color='#00000022') +
  ggtitle('Literacy versus Age Grouped by Subject')

pdf('literacy_versus_age.pdf', width=6, height=3.75)
p
dev.off()
p

```


png: 2



```
In [5]: p <- ggplot(literacy.data, aes(x=literacy, fill=follow.per)) +
  geom_histogram(bins=8) +
  scale_fill_brewer('Follow Period', palette='Set1', guide=FALSE) +
  facet_wrap(~follow.per, ncol=1) +
  ggtitle('Literacy Distribution by Follow Period')

pdf('literacy_by_follow_period.pdf', width=6, height=6)
p
dev.off()
p
```

png: 2



Simulation Study: Random versus Fixed Effect Intercept

```
In [6]: lmer.simulation.samples <- do.call(
  rbind, mclapply(replicate(8192, literacy.covariates.data, simplify=FALSE),
    function(covariates) {
      literacy <- make.response(covariates)
      model <- lmer(literacy ~ delta.age + (1|subject), data=covariates, REML=TRUE)
      var.corr <- VarCorr(model)
      data.frame(
        delta.age=fixef(model)[['delta.age']],
        delta.age.std.error=summary(model)$coefficients['delta.age', 'Std. Error'],
        sigma=attributes(var.corr)$sc,
        sigma.random=attributes(var.corr$subject)$stddev[['(Intercept)']]
      ), mc.cores=4))
```

```
In [7]: lm.simulation.samples <- do.call(
  rbind, mclapply(replicate(8192, literacy.covariates.data, simplify=FALSE),
    function(covariates) {
      literacy <- make.response(covariates)
      model <- lm(literacy ~ subject + delta.age, data=covariates)
      model.summary <- summary(model)
      data.frame(
        delta.age=model$coefficients[['delta.age']],
        delta.age.std.error=model.summary$coefficients['delta.age', 'Std. Error'],
        sigma=model.summary$sigma)
    }, mc.cores=4))
```

```
In [8]: summarize.simulation <- function(data) {
  l <- list(
    `\$\\mathbb{E}\\left[\\hat{\\beta}_L\\right]$`=mean(data$delta.age),
    `\$\\mathbb{E}\\left[\\hat{\\sigma}_L\\right]$`=mean(data$delta.age.std.
error),
    `Sample $\\hat{\\sigma}_L$`=sd(data$delta.age),
    `\$\\mathbb{E}\\left[\\hat{\\sigma}\\right]$`=mean(data$sigma))
  l[['\$\\mathbb{E}\\left[\\hat{\\sigma}_L\\right]$']] <- if ('sigma.random' %in% name
s(data)) {
    mean(data$sigma.random)
  } else {
    NA
  }
  l
}
```

```
In [9]: simulation.comparison <- as.data.frame(rbind(
  data.table(lmer.simulation.samples[, summarize.simulation(.SD)]),
  data.table(lm.simulation.samples[, summarize.simulation(.SD)]))
row.names(simulation.comparison) <- c('Random Effects Intercept', 'Fixed Effects Intercept')
simulation.comparison
```

	$\mathbb{E} \left[\hat{\beta}_L \right]$	$\mathbb{E} \left[\hat{\sigma}_{\hat{\beta}_L} \right]$	Sample $\hat{\sigma}_{\hat{\beta}_L}$	$\mathbb{E} \left[\hat{\sigma} \right]$	$\mathbb{E} \left[\hat{\sigma}_\gamma \right]$
Random Effects Intercept	1.2779277	0.3208065	0.3279972	10.007446	24.3412
Fixed Effects Intercept	0.9969691	0.3213570	0.3199906	9.980246	NA

```
In [10]: print(xtable(simulation.comparison,
  caption=paste('\small Results of a simulation study comparing modeling the ',
    'subject-specific intercepts as a random effect or fixed effect.',
    'Parameter estimates were averaged over simulations. Standard error
s',
    'for $\\hat{\\beta}_L$ are calculated two ways:',
    '(1) assuming the model is correct',
    '($\\mathbb{E}\\left[\\hat{\\sigma}_L\\right]$)',
    'and (2) using the $\\hat{\\beta}_L$ samples',
    '(Sample $\\hat{\\sigma}_L$).'),
  label='tab:simulation_comparison',
  digits=c(0, 6, 6, 6, 6, 6)),
  booktabs=TRUE,
  sanitize.colnames.function=identity,
  sanitize.row.names.function=identity,
  size='small',
  file='simulation_comparison.tex')
```

σ_γ Calculation

```
In [11]: base.literacy <- make.base.literacy(literacy.covariates.data[follow.per==0]$base.age)
variance.random.intercept <- mean((base.literacy - mean(base.literacy))^2)
sqrt(variance.random.intercept)
```

24.4330554781836

Expected $\hat{\beta}$

GLS

```
In [12]: expect.beta.hat.gls <- function(residual.variance) {
  subject.covariance.inv <- chol2inv(chol(
    as.matrix(nlme::pdCompSymm(variance.random.intercept + diag(3)*residual.variance))))
  projected.data <- lapply(levels(literacy.covariates.data$subject), function(i) {
    data <- literacy.covariates.data[J(i)]
    X <- cbind(1, data$delta.age)
    tX <- t(X)
    list(X=tX %*% subject.covariance.inv %*% X,
         y=tX %*% subject.covariance.inv %*% make.response(data, 0))
  })
  projected.X <- Reduce(
    `+`, lapply(projected.data, function(projection) { projection$X }))
  projected.y <- Reduce(
    `+`, lapply(projected.data, function(projection) { projection$y }))
  r <- chol(projected.X) # Upper triangular
  beta.hat <- as.vector(backsolve(r, forwardsolve(r, projected.y, upper.tri=TRUE, transpose=TRUE)))
  names(beta.hat) <- c('(Intercept)', 'delta.age')
  beta.hat
}
```

Ridge Regression

```
In [13]: expect.beta.hat.ridge <- function(residual.variance) {
  X <- cbind(1, model.matrix(~ 0 + delta.age + subject, literacy.covariates.data))
  y <- make.response(literacy.covariates.data, 0)
  # Penalize only random effects.
  Q <- diag(ncol(X))
  Q[1, 1] <- 0
  Q[2, 2] <- 0
  # Estimate beta
  tX <- t(X)
  r <- chol(tX %*% X + residual.variance/variance.random.intercept*Q)
  beta.hat <- as.vector(backsolve(r, forwardsolve(r, tX %*% y, upper.tri=TRUE, transpose=TRUE)))
  names(beta.hat) <- colnames(X)
  names(beta.hat)[1] <- '(Intercept)'
  beta.hat
}
```

Penalized Regression

```
In [14]: expect.beta.hat.penalized <- function(residual.variance) {
  coef(penalized(make.response(literacy.covariates.data, 0),
    unpenalized=~delta.age,
    penalized=~subject,
    data=literacy.covariates.data,
    lambda2=residual.variance/variance.random.intercept))
}
```

```
In [15]: expect.beta.hat.gls(100)['delta.age']
expect.beta.hat.ridge(100)['delta.age']
expect.beta.hat.penalized(100)['delta.age']
```

delta.age: 1.27288323972858

delta.age: 1.27288323972858

delta.age: 1.27288323972857

Bias as a Function of σ^2

```
In [16]: bias.data <- data.table(residual.variance=c(seq(0.1, by=0.1, length.out=9), 1:1000))
bias.data[,bias:=do.call(c, mclapply(
  residual.variance,
  function(residual.variance) { expect.beta.hat.ridge(residual.variance)[['delta.age']] },
  mc.cores=4)) - 1]

p <- ggplot(bias.data, aes(x=residual.variance, y=bias)) +
  geom_line() +
  scale_x_continuous(expression(sigma^2), breaks=seq(0, 1000, by=100)) +
  scale_y_continuous(expression(Bias(hat(beta)[L])), breaks=seq(0, 1.75, by=0.25)) +
  ggtitle('Bias as a Function of Residual Variance')

pdf('bias_variance_plot.pdf', width=6, height=3.75)
p
dev.off()
p
```

png: 2

